

A Fast Restart Mechanism for Checkpointing in Networked Environments

Yawei Li and [Zhiling Lan](#)

Department of Computer Science
Illinois Institute of Technology

Email: {liyawei,lan}@iit.edu

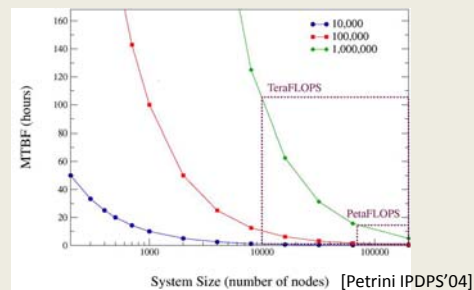
Outline

- Motivation
- FREM: a Fast Restart Mechanism
 - Design & implementation
- Experiments
- Summary



Reliability Crisis

- Failures are inevitable in large-scale systems
 - MTBF drops dramatically in large systems



- Long-running programs are forced to restart more frequently

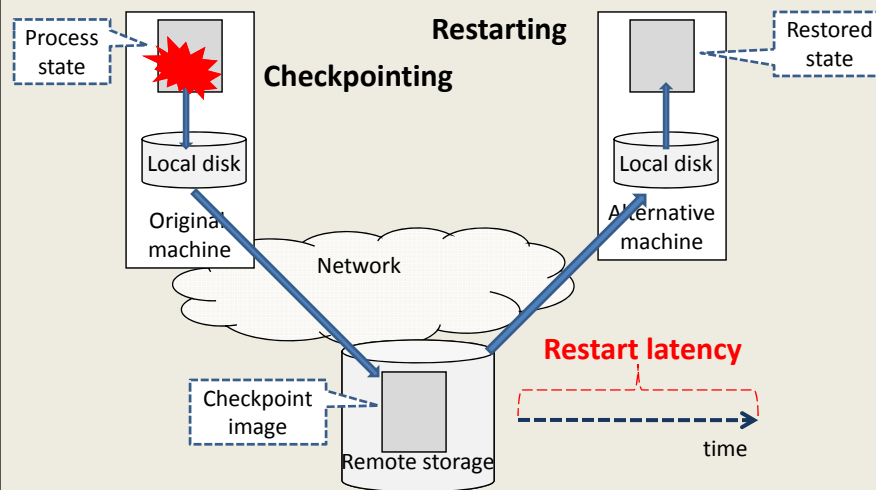


Checkpoint/Recovery

- C/R saves program state onto stable storage
 - Reduce rollback cost after restart
 - In networked systems, enable remote process restart on alternative resource
- **C/R is not a silver bullet!**
 - Extensive studies on reducing checkpointing overhead & determining an optimal frequency
 - Little work on optimizing its restart



An Illustrative Example



Problem Statement

- Currently, a restart requires the entire checkpoint image before it can proceed
 - Substantial restart latency in networked environments
 - Network transmission and I/O operation time
- Insatiable data demand from applications leads to larger checkpoint size
 - Thus, longer restart latency



The Pain of Restart Latency

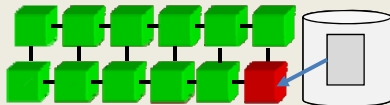
- Restart latency directly contributes to program downtime
 - $Avail = (up\ time) / (up\ time + down\ time)$

Amazon Did NOT Just Lose \$3 Mil of Revenue In 90 Mins Downtime (AMZN)

Henry Blodget | June 6, 2008 5:19 PM

Around 1:30 ET, Amazon's site crashed (AMZN). Around 3:30, we got an email with what will no doubt be the first of many estimates that try to quantify how much revenue Amazon just "lost" (The first estimate was \$1.8 million an hour, or about \$3 million during the 90 minutes of downtime.)

But how much revenue did Amazon really "lose" during this outage? Most likely only a small fraction of this.



Parallel jobs make it worse:
all processes stall during restart!



ILLINOIS INSTITUTE OF TECHNOLOGY

Zhiling Lan

7

FREM: A Fast REstart Mechanism

- Provide a quick restart through *latency hiding*
 - Overlapping process execution with checkpoint image retrieval
- Based on two key observations
 - Data locality: only a small portion of the checkpoint image is immediately needed after resurrection
 - Predictable pattern: data access patterns of restart from checkpoint can be known



ILLINOIS INSTITUTE OF TECHNOLOGY

Zhiling Lan

8

Main Idea

1) Post-checkpoint tracking

- Record the *touch set* following each CKP within the *tracking window*
 - Touch set: defined as the intersection of the process address space saved in the checkpoint and its working set
 - It shows which data is immediately needed for restart from the image file

2) Fast restarting

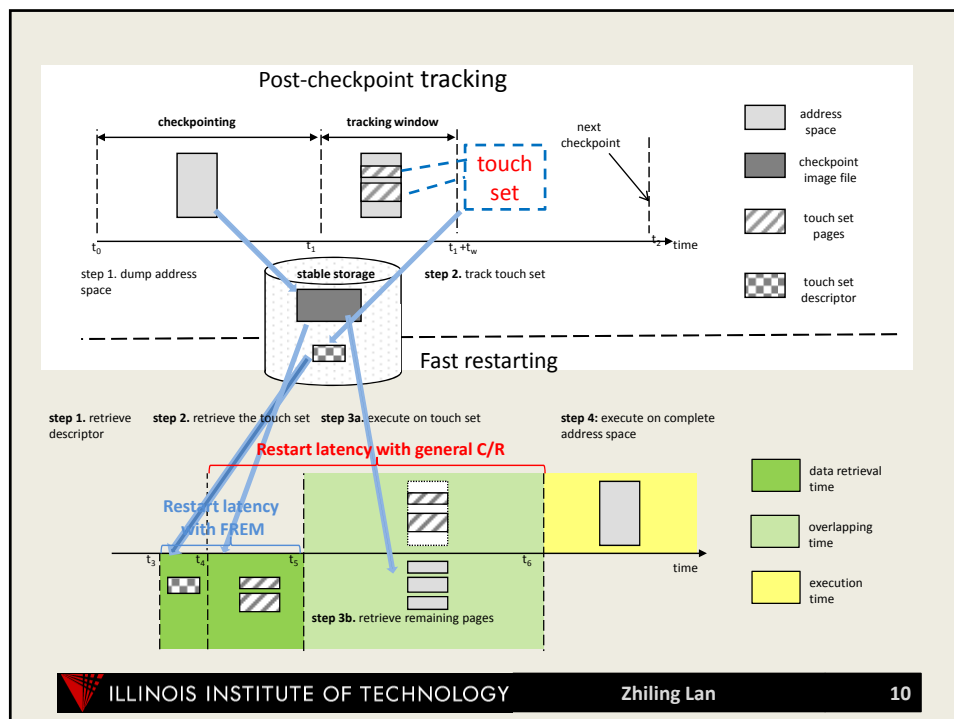
- Restart the process as soon as the touch set is retrieved
- Overlap the process execution with data retrieval of the remaining image



ILLINOIS INSTITUTE OF TECHNOLOGY

Zhiling Lan

9



ILLINOIS INSTITUTE OF TECHNOLOGY

Zhiling Lan

10

Technical Challenges

- 1) How to accurately identify the touch set?
- 2) How to appropriately set the tracking window?
- 3) How to effectively load the partial image?



Technical Challenges

- 1) How to accurately identify the touch set?
- 2) How to appropriately set the tracking window?
- 3) How to effectively load the partial image?



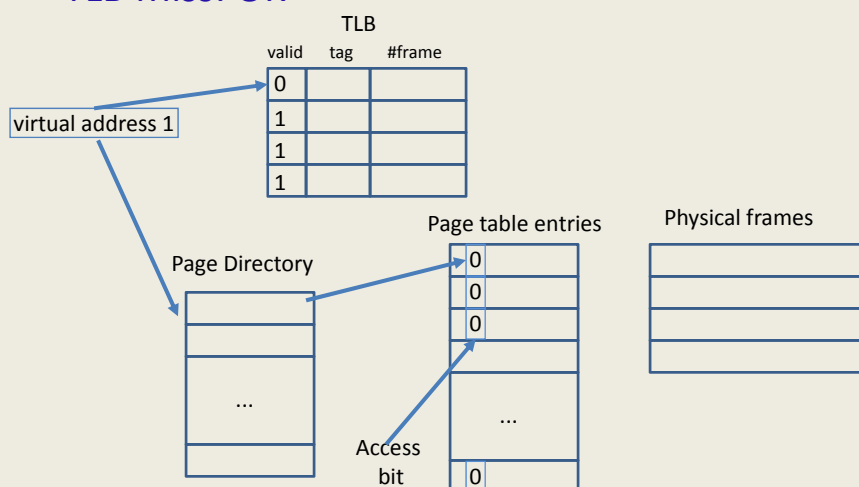
Identifying Touch Set

- Find out the data that is immediately needed for restart from the image file
- Two types of errors
 - False negative – a page of interest is missed
 - False positive – a page not interested is included
- Sources of inaccuracy:
 - Hardware bypassing, page swapping, dynamic memory usage



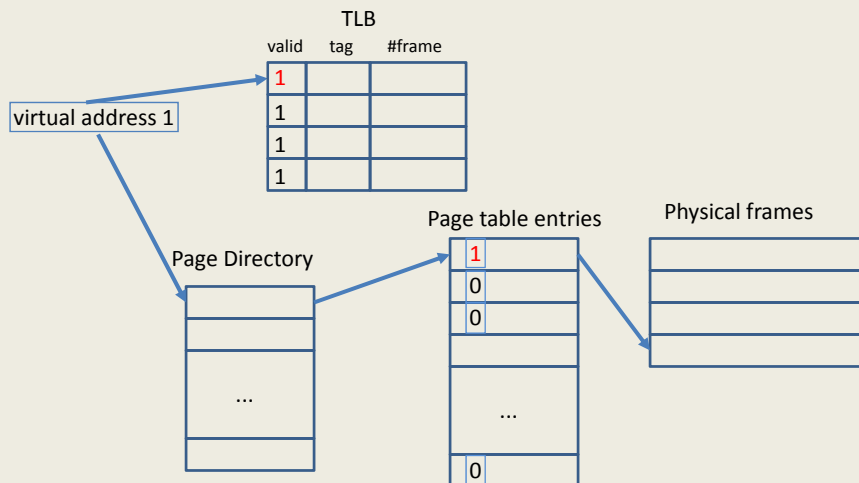
Hardware Bypassing – TLB

- TLB Miss: OK



Hardware Bypassing – TLB

■ TLB Miss: OK



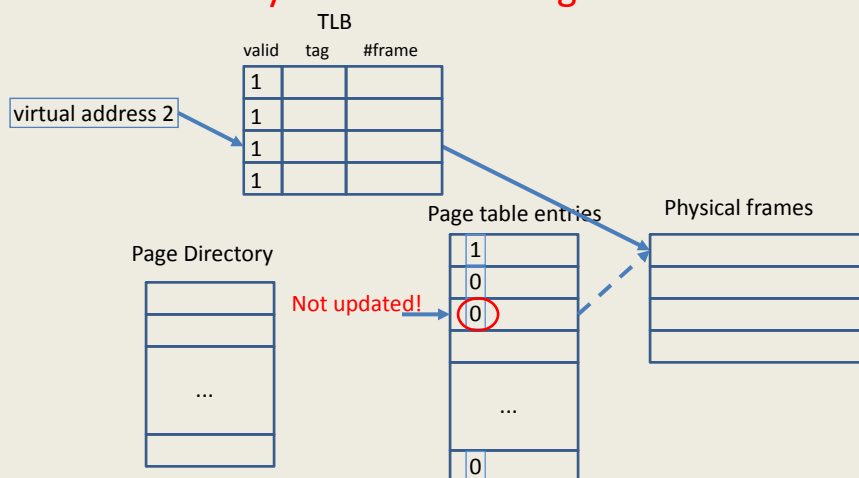
ILLINOIS INSTITUTE OF TECHNOLOGY

Zhiling Lan

15

Hardware Bypassing – TLB

■ TLB Hit: May lead to false negatives



ILLINOIS INSTITUTE OF TECHNOLOGY

Zhiling Lan

16

Hardware Bypassing – TLB

- Solution: invalidate TLB entry to enforce TLB miss

| TLB | | |
|-------|-----|--------|
| valid | tag | #frame |
| 0 | | |
| 1 | | |
| 1 | | |
| 1 | | |

Page Directory

| |
|-----|
| |
| |
| ... |
| |

Step 1: reset access bits in PTE when tracking starts

Page table entries

| |
|-----|
| 0 |
| 0 |
| 0 |
| ... |
| 0 |

Physical frames

| |
|--|
| |
| |
| |
| |



ILLINOIS INSTITUTE OF TECHNOLOGY

Zhiling Lan

17

Hardware Bypassing – TLB

- Solution: invalidate TLB entry to enforce TLB miss

| TLB | | |
|-------|-----|--------|
| valid | tag | #frame |
| 0 | | |
| 1 | | |
| 0 | | |
| 1 | | |

Page Directory

| |
|-----|
| |
| |
| ... |
| |

Step 2: invalidate corresponding TLB entries immediately

Page table entries

| |
|-----|
| 0 |
| 0 |
| 0 |
| ... |
| 0 |

Physical frames

| |
|--|
| |
| |
| |
| |



ILLINOIS INSTITUTE OF TECHNOLOGY

Zhiling Lan

18

Hardware Bypassing – DMA

- Hardware devices access system memory for reading/writing independently of the CPU
 - Direct I/O, zero copy optimization
- Solution: include the mapped DMA pages in the touch set



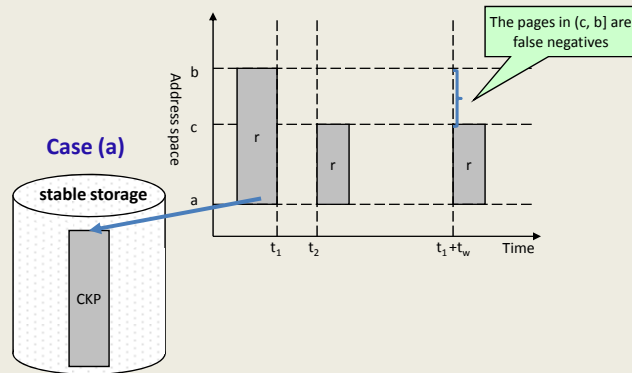
Page Swapping

- OS page swapping may cause false negatives
 - Page swapping algorithm clears the access bits when do the page table walk
- Solution: intercept *kswapd* to record the page per swapping



Dynamic Memory Usage

- Complications incurred by dynamic memory usage



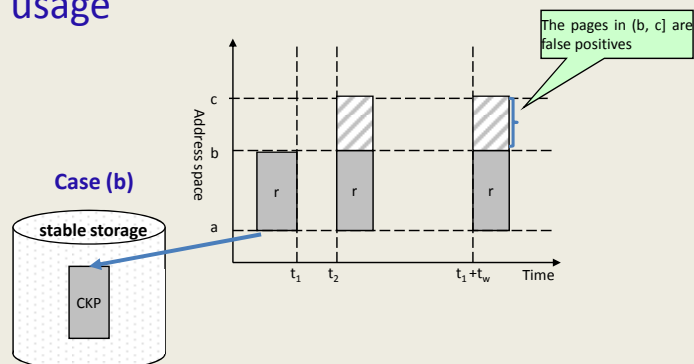
ILLINOIS INSTITUTE OF TECHNOLOGY

Zhiling Lan

21

Dynamic Memory Usage

- Complications incurred by dynamic memory usage



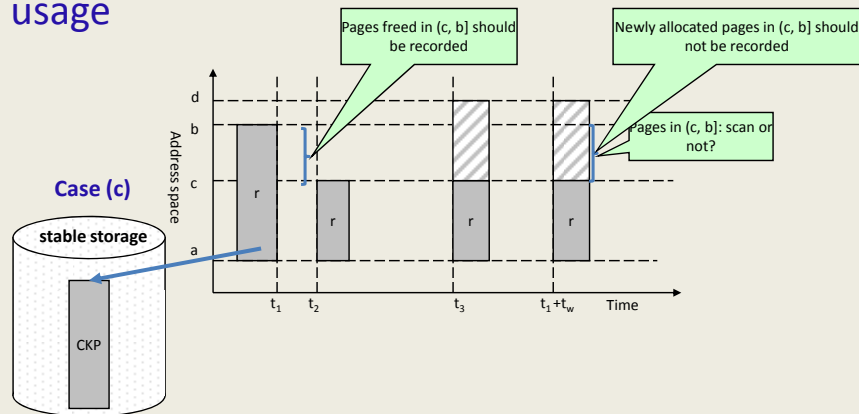
ILLINOIS INSTITUTE OF TECHNOLOGY

Zhiling Lan

22

Dynamic Memory Usage

- Complications incurred by dynamic memory usage



ILLINOIS INSTITUTE OF TECHNOLOGY

Zhiling Lan

23

Dynamic Memory Usage

- Observation:
 - The touch set is always a subset of the checkpoint image, which monotonically decreases
- Solution
 - Record the pages saved in the checkpoint image as the *candidate pages*;
 - Intercept *do_munmap* to capture memory shrinking
 - Scan the intersection of the *candidate pages* and the pages to be freed
 - Update the *candidate pages* to exclude the intersection



ILLINOIS INSTITUTE OF TECHNOLOGY

Zhiling Lan

24

Experiments

- A prototype implementation with the BLCR tool in Linux 2.6.22
 - Use a red-black tree for fast touch set tracking
 - Instrument the kernel in a non-intrusive way
- FREM-enhanced BLCR vs. BLCR
 - Evaluated with the SPEC CPU2006 suite
 - Two network settings

| Network | Network Latency | Network Bandwidth |
|---------|-----------------|-------------------|
| SLOW | 200 ms | 7MB/s |
| FAST | 70 ms | 32MB/s |



Checkpoint image size

Restart latency

| Application (input set) | W (MB) | Touch set size (MB) | | RL with BLCR (s) | | RL with FREM (s) | |
|----------------------------|-------------|------------------------|------|-----------------------|-------|--------------------|---------------|
| | | FAST | SLOW | FAST | SLOW | FAST | SLOW |
| 1: astar (1) | 280 | 34 | 44 | 49.4 | 80.1 | 6.0 (87.88%) | 12.6 (84.26%) |
| 2: bzip2 (5) | 847 | 45 | 152 | 161.3 | 254.1 | 8.6 (94.64%) | 45.6 (82.04%) |
| 3: bzip2 (6) | 609 | 64 | 244 | 109.3 | 176.0 | 11.5 (89.49%) | 70.5 (59.97%) |
| 4: dealII | 239 | 12 | 28 | 31.0 | 57.2 | 1.6 (94.80%) | 6.8 (88.13%) |
| 5: gamess (1) | 629 | 5 | 9 | 112.0 | 180.9 | 0.8 (99.28%) | 2.6 (98.57%) |
| 6: gcc (4) | 311 | 48 | 82 | 53.4 | 87.6 | 11.6 (78.20%) | 23.0 (73.75%) |
| 7: gcc (6) | 771 | 216 | 211 | 136.7 | 221.2 | 38.2 (72.05%) | 60.6 (72.62%) |
| 8: lbm | 409 | 402 | 402 | 73.6 | 118.5 | 72.3 (1.80%) | 116.3 (1.80%) |
| 9: mcf | 839 | 394 | 827 | 151.0 | 242.8 | 70.9 (53.03%) | 239.5 (1.37%) |
| 10: perl (1) | 171 | 31 | 50 | 24.7 | 43.5 | 4.4 (82.16%) | 12.8 (70.66%) |
| 11: soplex (2) | 490 | 186 | 191 | 89.3 | 142.9 | 33.9 (62.05%) | 55.7 (61.06%) |
| 12: wrf | 685 | 37 | 346 | 117.8 | 192.9 | 54.5 (53.78%) | 97.9 (49.25%) |



Runtime Overhead during Post-Checkpoint Tracking (in milliseconds)

| Application (input set) | Scan time | Search and insertion time | Descriptor I/O time | Tracking overhead |
|-----------------------------|-----------|------------------------------|------------------------|----------------------|
| 1: astar (1) | 18.747 | 13.517 | 1.360 | 33.632 |
| 2: bzip2 (5) | 48.790 | 0.257 | 0.184 | 49.231 |
| 3: bzip2 (6) | 38.763 | 0.448 | 0.149 | 39.360 |
| 4: dealII | 4.663 | 1.718 | 0.220 | 6.601 |
| 5: gamese (1) | 36.711 | 1.678 | 0.298 | 38.687 |
| 6: gcc (4) | 29.297 | 4.641 | 0.549 | 34.487 |
| 7: gcc (6) | 43.734 | 14.954 | 1.488 | 60.176 |
| 8: lbm | 21.832 | 0.462 | 0.119 | 22.413 |
| 9: mcf | 3.390 | 0.175 | 0.161 | 34.234 |
| 10: perl (1) | 18.573 | 13.630 | 1.350 | 33.553 |
| 11: soplex (2) | 29.749 | 14.359 | 1.459 | 45.566 |
| 12: wrf | 41.594 | 4.643 | 0.608 | 46.845 |



ILLINOIS INSTITUTE OF TECHNOLOGY

Zhiling Lan

27

Runtime Overhead during Fast Restart (in seconds)

| Application (input set) | Remaining image (MB) | Duration of overlapping | Fast restart overhead |
|-----------------------------|-------------------------|----------------------------|--------------------------|
| 1: astar (1) | 162 | 19 | 7.1 |
| 2: bzip2 (5) | 476 | 55 | 13.1 |
| 3: bzip2 (6) | 250 | 48 | 11.8 |
| 4: dealII | 144 | 14 | 10.2 |
| 5: gamese (1) | 424 | 60 | 21.7 |
| 6: gcc (4) | 157 | 19 | 10.6 |
| 7: gcc (6) | 560 | 77 | 22.7 |
| 8: lbm | 5 | 0.8 | 0.1 |
| 9: mcf | 8 | 1.4 | 0.3 |
| 10: perl (1) | 83 | 12 | 6.9 |
| 11: soplex (2) | 205 | 30 | 4.8 |
| 12: wrf | 231 | 35 | 7.8 |



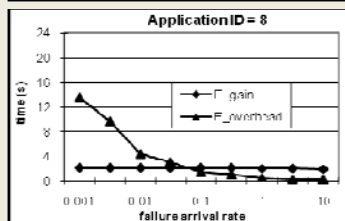
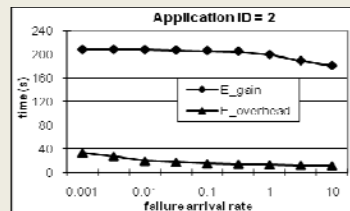
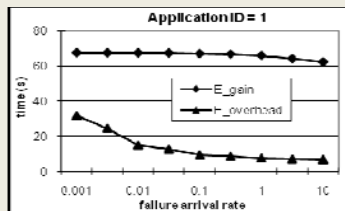
ILLINOIS INSTITUTE OF TECHNOLOGY

Zhiling Lan

28

Performance Evaluation

- statistically evaluate application performance **between two restarts**



E_{gain} = expected restart improvement
between 2 restarts;
 E_{overhead} = expected runtime overhead
introduced between 2 restarts



ILLINOIS INSTITUTE OF TECHNOLOGY

Zhiling Lan

29

Related Work

- Existing work on fast restart
 - OS, Database, Internet services, ...
 - The novelty of FREM is its ability to reduce restart latency by recording data accesses after CKP*
- Traditional work on C/R
 - Focus on checkpoint optimization, e.g., determining an optimal checkpoint frequency
 - FREM complements these studies*
- Demand paging
 - FREM selectively restores the pages needed by tracking data access patterns*



ILLINOIS INSTITUTE OF TECHNOLOGY

Zhiling Lan

30

Summary

- Have presented a novel mechanism for fast recovery of C/R applications
 - Through a user-transparent system support
- Have demonstrated its effectiveness
 - FREM can reduce process restart latency by 61.96% on average
- Future work
 - Adaptive tracking window estimation
 - Better image loading mechanism
 - Integration with checkpointing tools



Questions?

FENCE Project Website:

<http://www.cs.iit.edu/~zlan/fence.html>